

Computer Science I

Arrays

Dr. Chris Bourke
cbourke@cse.unl.edu

Outline

1. Introduction
2. Using Arrays
3. Dynamic Arrays
4. Memory Management
5. Arrays & Functions
6. Multidimensional Arrays
7. Shallow vs. Deep Copies

Part I: Introduction

Arrays

- ▶ Rarely do we deal with only one piece of data
- ▶ Usually more than one number, string, object, etc. must be stored and processed
- ▶ *Collections* of data can be stored in *arrays*
- ▶ An “array” is an *ordered series or arrangement*

Arrays

In code:

- ▶ Arrays are collections of ordered data stored *contiguously* in memory
- ▶ *ordered* is not the same as *sorted*
- ▶ Have a single identifier (name)
- ▶ Size is *fixed* when created
- ▶ You access individual elements in an array with an *index*
- ▶ Arrays are 0-indexed: first element is at index 0, the second at index 1, etc.
- ▶ An array of size n has the last element at index $n - 1$
- ▶ Indexing is usually done with the square brackets `[]`

Example

index	0	1	2	3	4	5	6	7	8
contents	48	9	17	5	29	72	42	101	32


Indexing

- ▶ Suppose `arr` is an integer (4 bytes each) array
- ▶ `arr` is actually a *memory address*
- ▶ i -th element is at `arr[i]`
- ▶ Indexing automatically computes a *memory offset*
- ▶ i -th element is

$$i \times 4$$

bytes away from the beginning of the array

Memory Offsets



Part II: Using Arrays

Using Arrays

- ▶ Static arrays are allocated on the program stack
- ▶ Declaration specifies size

```
1 int arr[10];  
2 double numbers[20];
```

- ▶ Once declared, indexing can be used to access values

```
1 arr[0] = 42;  
2 arr[1] = 12;  
3 arr[2] = arr[0] + 20;  
4 arr[9] = 3.75; //truncation  
5  
6 printf("a[0] = %d\n", a[0]);
```

Alternative Syntax

Declaration/initialization

- ▶ You can declare and initialize an array at the same time

```
1 int primes[] = { 2, 3, 5, 7, 11, 13, 17 };
```

- ▶ Size specification is *optional*
- ▶ Example of *code elision*
- ▶ Problem: You still need to keep track of the size of the array

Variable Length Arrays

- ▶ C99+ allows you to declare an array with a variable size

```
1 int n = 10;  
2 int arr[n];
```

- ▶ Just because you *can* (or more accurately *might*) be able to do this, doesn't mean you should
- ▶ Avoid in general

Pitfalls

Uninitialized Arrays

- ▶ Like regular variables, there is *no default value*
- ▶ `arr[3]` was not set, its value could be anything
- ▶ Never make assumptions about uninitialized variables
- ▶ Always initialize yourself

Pitfalls

Out-of-bounds indexing

- ▶ Accessing invalid indices is *undefined behavior*

```
1  int arr[10];
2  ...
3  arr[10] = 42;
4  arr[-1] = 21;
```

- ▶ May lead to:
 - ▶ A segmentation fault, bus error
 - ▶ Corrupted memory
 - ▶ Incorrect results
- ▶ *Your responsibility to do bookkeeping*

Pitfalls

Book Keeping

- ▶ You must always keep track of the size of an array
- ▶ In general there is no way to determine the size of an array
- ▶ Only in very limited situations (static arrays)
- ▶ Arrays should be accompanied by an integer variable to keep track of its size
- ▶ Idiomatic loops over arrays
- ▶ Demonstration

Pitfalls

Book Keeping

```
1  int n = 10;
2  int primes[] = { 2, 3, 5, 7, 11, 13, 17, 19, 23, 29 };
3  int sum = 0;
4
5  for(int i=0; i<n; i++) {
6      sum += primes[i];
7  }
```

Part III: Dynamic Arrays

Static Arrays are Insufficient

- ▶ Static arrays are allocated on the program stack
- ▶ Inside a stack frame in which it is declared
- ▶ Stack space is *limited*
- ▶ 8MB (large) to 64k or even 8k (embedded systems)
- ▶ Demonstration

Static Arrays are Insufficient

- ▶ Stack is small, inappropriate to hold even “moderately” sized arrays
- ▶ Other disadvantages: static arrays cannot be returned from functions
- ▶ Best to not use static arrays at all
- ▶ Don’t abuse the stack space, it is small and defenseless
- ▶ Better solution: use *dynamic arrays*

Dynamic Memory & Arrays

- ▶ Dynamic memory is allocated in a program's *heap*
- ▶ Stack: highly organized, efficient, but small/limited
- ▶ Heap: Less organized, less efficient, but much larger
- ▶ Dynamically allocate memory on the heap using `malloc()` (**Memory Allocation**)

malloc

- ▶ Located in the standard library `stdlib.h`
- ▶ Takes one argument: the number of bytes you want to allocate
- ▶ Use `sizeof()` to determine how many bytes each type of variable takes
- ▶ Returns a generic *void pointer*: `void *`
- ▶ A void pointer points to a generic memory location that can be *cast* to any type you want
- ▶ Returns `NULL` if unsuccessful
- ▶ Demonstration

Part IV: Memory Management

Overview

- ▶ Memory on the stack is “cleaned up” when stack frames are removed
- ▶ Memory on the heap is *not* automatically cleaned up when it is no longer needed
- ▶ It is *your* responsibility to “clean up” dynamically allocated memory when you no longer need it
- ▶ Failure to do so or failure to do so correctly can lead to:
 - ▶ Memory Leaks
 - ▶ Reduced performance
 - ▶ Illegal memory access/segmentation faults

Freeing Memory

- ▶ To clean up memory you “free” it
- ▶ Standard library function:
`void free(void *)`
- ▶ Takes a single argument: a pointer to dynamically allocated memory
- ▶ Demonstration

Demonstration

```
1  #include<stdlib.h>
2  #include<stdio.h>
3  int main(int argc, char **argv) {
4
5      int n = 100;
6      int *arr = (int *) malloc(n * sizeof(int));
7
8      //process the array
9      for(int i=0; i<n; i++) {
10         arr[i] = (i+1);
11     }
12
13     free(arr);
14
15     return 0;
16 }
```

Pitfall

Dangling Pointers

- ▶ Basic usage is easy, though there are many pitfalls
- ▶ Once freed, dynamic memory cannot/should not be accessed
- ▶ Best to *reset* the pointer to `NULL`
- ▶ Otherwise, it is a *dangling pointer*
- ▶ Demonstration

Pitfall

Dangling Pointers

```
1  int *arr = (int *) malloc(n * sizeof(int));
2  //...
3  free(arr);
4
5  //arr still points to a memory location, but is no longer valid
6  printf("arr points to %p\n", arr);
7  //accessing is undefined behavior:
8  arr[0] = 42;
9
10 //best to reset to NULL;
11 free(arr);
12 arr = NULL;
```

Pitfall

Memory Ownership

- ▶ Different sections of code may "own" memory and be responsible for its management
- ▶ Stack frames are "owned" by the program/function: the program is responsible for clean up
- ▶ Ownership may be transferred: `malloc` transfers ownership to the calling function
- ▶ Ownership is a design issue/decision
- ▶ In general: only `free` memory if you own it
- ▶ Don't `free` memory before you are done with it
- ▶ Don't `free` freed memory

Pitfall

Memory Leaks

- ▶ Failure to properly clean up memory can lead to *memory leaks*
- ▶ A program holds on to memory it doesn't use
- ▶ Or: references are lost to memory that cannot be freed
- ▶ Program takes more and more resources
- ▶ Performance degrades, taking the entire system down with it
- ▶ Demonstration

Part V: Arrays & Functions

Using Arrays With Functions

- ▶ Goal: use arrays with functions
- ▶ Pass arrays to functions
- ▶ Return arrays from functions
- ▶ Recall: you always need to do your own *bookkeeping* with arrays
- ▶ Anytime you pass an array, you also need to pass its *size*
- ▶ Anytime you return an array, you need a way to implicitly determine its size

Demonstration

Write a function that takes an array of integers and returns the sum of its values.

Demonstration

```
1  /**
2   * This function takes an integer array (of size n) and
3   * returns the sum of its elements. It returns 0 if the
4   * array is NULL.
5   */
6  int sum(int *arr, int n) {
7
8      if(arr == NULL) {
9          return 0;
10     }
11     int total = 0;
12     for(int i=0; i<n; i++) {
13         total += arr[i];
14     }
15     return total;
16 }
```

The const Keyword

- ▶ Arrays are *always* passed by reference in C
- ▶ It is possible to make changes to their contents
- ▶ We don't always want this
- ▶ We can add the keyword `const` to prevent changes to the array
- ▶ The compiler checks for changes and generates an error if this "promise" is violated
- ▶ Design issue, not a full guarantee
- ▶ Demonstration

Returning Arrays

- ▶ Functions can create and return arrays
- ▶ You *cannot* return static arrays
- ▶ Only dynamic arrays can be returned
- ▶ Demonstration

Returning Arrays

```
1  #include <stdlib.h>
2  #include <stdio.h>
3
4  int * foo() {
5      int b[3];
6      b[0] = 10;
7      b[1] = 20;
8      b[2] = 30;
9      return b;
10 }
11
12 int main(int argc, char **argv) {
13     int *a = foo();
14     for(int i=0; i<3; i++) {
15         printf("a[%d] = %d\n", i, a[i]);
16     }
17     return 0;
18 }
```

Exercises

- ▶ Write a function that takes an integer n and returns an array of n integers, all initialized to 1
- ▶ Write a function that takes an integer array and returns a new *copy* of the array with all instances of zero removed

Part VI: Multidimensional Arrays

Multidimensional Arrays

- ▶ Up to now: 1-dimensional arrays
- ▶ You can have arrays with more than one dimension
- ▶ 2-D arrays:
 - ▶ Rows & columns
 - ▶ Tabular data
 - ▶ Matrices
- ▶ 3-D arrays:
 - ▶ Rows, columns, & “lanes”
 - ▶ 3-dimensional data
- ▶ 4+ dimensional arrays:
 - ▶ Rethink what you're doing
- ▶ Focus on 2-D arrays

Multidimensional Arrays


- ▶ A pointer, `int *arr` points to a 1-dimensional array
- ▶ A “double” pointer, `int **mat` points to a “2-dimensional array”
- ▶ Technically points to an array of pointers
- ▶ Each pointer in the array points to an array of integers
- ▶ Demonstration

Multidimensional Arrays

```
1 int n = 10;
2 int **matrix = NULL;
3 matrix = (int **) malloc(n * sizeof(int*));
4 for(int i=0; i<n; i++) {
5     matrix[i] = (int *) malloc(n * sizeof(int));
6 }
```

Multidimensional Arrays

```
1 int n = 10;
2 int **matrix = NULL;
3 matrix = (int **) malloc(n * sizeof(int*));
4 for(int i=0; i<n; i++) {
5     matrix[i] = (int *) malloc(n * sizeof(int));
6 }
```



Usage

- ▶ Once created, you can access elements using two indices
- ▶ Usually: row-column interpretation
- ▶ `matrix[i][j]` accesses the i -th row and j -th column
- ▶ Use two nested for-loops to iterate over each row/column

```
1 for(int i=0; i<n; i++) {
2     for(int j=0; j<n; j++) {
3         matrix[i][j] = (2*i+3*j);
4     }
5 }
6 printf("last row/column value = %d\n", matrix[n-1][n-1]);
```

Demo

```
1 #include<stdlib.h>
2 #include<stdio.h>
3
4 int main(int argc, char **argv) {
5
6     int n = 3;
7     int **matrix = NULL;
8     matrix = (int **) malloc(n * sizeof(int*));
9     for(int i=0; i<n; i++) {
10         matrix[i] = (int *) malloc(n * sizeof(int));
11     }
12
13     int value = 1;
14     for(int i=0; i<n; i++) {
15         for(int j=0; j<n; j++) {
16             matrix[i][j] = value;
17             value++;
18         }
19     }
```

Clean Up

- ▶ Must do proper cleanup when freeing 2-D arrays
- ▶ Cannot simply free the matrix: `free(matrix)`
- ▶ Results in a memory leak
- ▶ You must free each row before you free the array of pointers

```
1 for(int i=0; i<n; i++) {
2     free(matrix[i]);
3 }
4 free(matrix);
```

Alternative: Contiguous Allocation

```
1 #include<stdlib.h>
2
3 int main(int argc, char **argv) {
4
5     int n = 5, m = 3;
6     int **arr = (int **) malloc(sizeof(int *) * n);
7     arr[0] = (int *) malloc(sizeof(int) * (n * m));
8
9     for(int i=1; i<n; i++) {
10         arr[i] = (*arr + (m * i));
11     }
12
13     int value = 1;
14     for(int i=0; i<n; i++) {
15         for(int j=0; j<m; j++) {
16             arr[i][j] = value;
17             value++;
18         }
19     }
```


Part VII: Shallow vs. Deep Copies

Shallow Copy

Consider the following piece of code, what does it print?

```
1 //create an array containing {10, 20, 30, 40, 50}
2 int n = 5;
3 int *a = (int *) malloc(n * sizeof(int));
4 for(int i=0; i<n; i++) {
5     a[i] = (i+1)*10;
6 }
7
8 //let's make a "copy"
9 int *b = a;
10 b[0] = 42;
11
12 //what is in a[0]?
13 printf("a[0] = %d\n", a[0]);
```


Shallow Copy




Shallow Copy

- ▶ A shared reference is a *shallow copy*
- ▶ Multiple pointers refer to the same memory location/array
- ▶ Changes to one reference affect the other
- ▶ Not typically what we want with a "copy"

Deep Copy

- ▶ In contrast: a *deep copy* is when we have two *separate* arrays with the same *contents*
- ▶ Two *different* memory locations
- ▶ Changes to one do not affect the other
- ▶ Example
- ▶ Demo: write a deep copy function

Deep Copy



Deep Copy

```
1  /**
2   * This function creates a deep copy of the
3   * given array.
4   */
5  int * deepCopy(const int *a, int n) {
6      if(a == NULL || n < 0) {
7          return NULL;
8      }
9      int *copy = (int *) malloc(n * sizeof(int));
10     for(int i=0; i<n; i++) {
11         copy[i] = a[i];
12     }
13     return copy;
14 }
```