

## Computer Science I

### Searching & Sorting

Dr. Chris Bourke  
cbourke@cse.unl.edu

## Outline

1. Introduction & Linear Search
2. Binary Search
3. Sorting: Selection Sort
4. Sorting: Quick Sort
5. Sorting in Practice
6. Function Pointers
7. Searching & Sorting in C

## Part I: Introduction & Linear Search

## Introduction

- ▶ Processing data is a fundamental operation in Computer Science
- ▶ Two fundamental operations in processing data are *searching* and *sorting*
- ▶ Form the basis or preprocessing step of many algorithms
- ▶ Large variety of algorithms have been developed

## Searching

- ▶ Given a *collection of elements*  $A = \{a_1, a_2, \dots, a_n\}$  and a *key*  $k$ , find an element that "matches"  $k$
- ▶ Collection: haystack, key: needle

## Searching

Very general problems statement:

- ▶ Collection: arrays, sets, lists, etc.
- ▶ Elements: integers, strings, structures, etc.
- ▶ "matches": could be any criteria!
- ▶ Variations:
  - ▶ Find the first/last such element
  - ▶ Find all such elements
  - ▶ Find extremal elements
- ▶ What do you do for unsuccessful searches?

## Linear Search

Potential Solution: Linear Search

- ▶ Basic idea: iterate through each element
- ▶ For each element, apply the “matching” criteria
- ▶ Stop at the first match
- ▶ If no such element, return a “flag” value

## Linear Search

A potential C solution:

- ▶ Take an array of integers
- ▶ An integer key  $k$
- ▶ Find the first element equal to  $k$
- ▶ Return its index
- ▶ Unsuccessful search:  $-1$  as a flag value

## Linear Search

```
1  /**
2   * This function takes an array of integers
3   * and searches it for the given key, returning
4   * the index at which it finds it, or -1 if no
5   * such element exists.
6   */
7  int linearSearch(const int *arr, int n, int key) {
8
9      for(int i=0; i<n; i++) {
10         if(arr[i] == key) {
11             //you found your needle...
12             return i;
13         }
14     }
15     //the needle was not found
16     return -1;
17 }
```

## Linear Search: Observations

Solution works


- ▶ Solution works but is less than ideal
- ▶ It only applies to arrays of integers
- ▶ Search arrays of `double` or strings or `Student` structures, etc.: copy-pasta
- ▶ Different search criteria (search `Student` by NUID or name): yet another implementation
- ▶ Ultimate goal: one single “generic” searching (and sorting) solution that will work with arrays of any type of data
- ▶ Can we do better?

## Part II: Binary Search & Comparison

## Binary Search: Basic Idea

- ▶ Can we do better than linear search?
- ▶ Suppose that the array is *sorted*: how might we exploit that structure?
- ▶ Searching for an element  $k$
- ▶ Examine the middle element,  $m$ :
  - ▶ If  $m = k$ : success!
  - ▶ If  $k < m$ :  $k$  must lie in the left-half of the array
  - ▶ If  $m < k$ :  $k$  must lie in the right-half of the array

Illustration



Example

Search for  $k = 42$ :

$l = 0, \quad r = 10, \quad m = 5$

index	0	1	2	3	4	5	6	7	8	9	10
contents	-3	2	4	4	9	12	34	42	102	157	180

Example

$12 < 42 = k$ :

$l = 6, \quad r = 10, \quad m = 5$

index	0	1	2	3	4	5	6	7	8	9	10
contents	-3	2	4	4	9	12	34	42	102	157	180

Example

$l = 6, \quad r = 10, \quad m = 8$

index	0	1	2	3	4	5	6	7	8	9	10
contents	-3	2	4	4	9	12	34	42	102	157	180

Example

$42 < 102$ :

$l = 6, \quad r = 7, \quad m = 8$

index	0	1	2	3	4	5	6	7	8	9	10
contents	-3	2	4	4	9	12	34	42	102	157	180

Example

$34 < 42 = k$ :

$l = 6, \quad r = 7, \quad m = 6$

index	0	1	2	3	4	5	6	7	8	9	10
contents	-3	2	4	4	9	12	34	42	102	157	180

## Example

$$a_7 = 42 = k$$

$$l = 7, \quad r = 7, \quad m = 7$$

index	0	1	2	3	4	5	6	7	8	9	10
contents	-3	2	4	4	9	12	34	42	102	157	180

## Recursive Code

```
1 int binarySearch(const int *arr, int l, int r, int k) {
2     if(l > r) {
3         return -1;
4     } else {
5         int m = (l + r) / 2; //bad in practice
6
7         if(arr[m] == k) {
8             return m;
9         } else if(k < arr[m]) {
10            return binarySearch(arr, l, m-1, k);
11        } else if(arr[m] < k) {
12            return binarySearch(arr, m+1, r, k);
13        }
14    }
15 }
```

## Iterative Code

```
1 int binarySearch(const int *arr, int n, int k) {
2     int l = 0;
3     int r = n-1;
4     while(l <= r) {
5         int m = (l + r) / 2; //bad in practice
6
7         if(arr[m] == k) {
8             return m;
9         } else if(k < arr[m]) {
10            r = m - 1;
11        } else if(arr[m] < k) {
12            l = m+1;
13        }
14    }
15    return -1;
16 }
```

## Analysis

- ▶ Which is better? How much better?
- ▶ How much “work” does each algorithm perform?
- ▶ Suppose we search an array of  $n$  elements
- ▶ How many *comparisons* does each search perform?

## Linear Search Analysis

- ▶ Best case scenario: you get lucky and immediately find the element, making one single comparison
- ▶ Worst Case: you are unlucky and make all  $n$  comparisons
- ▶ Average case scenario:  $\approx \frac{n}{2}$  comparisons
- ▶ Called *linear search* because the work is *linearly* proportional to the array size

## Binary Search

- ▶ Worst case scenario: unsuccessful search
- ▶ Or: when the list size is cut down to size 1
- ▶ Each comparison cuts the array (roughly) in half
- ▶ After first iteration:  $\frac{n}{2}$
- ▶ After second:  $\frac{n}{4}$

### Binary Search

- ▶ After third:  $\frac{n}{8}$
- ▶ After  $k$  iterations:  $\frac{n}{2^k}$
- ▶ Stops when  $\frac{n}{2^k} = 1$
- ▶ Solve for  $k$ :  $k = \log_2(n)$
- ▶ Roughly only  $\log_2(n)$  comparisons are made.

### Comparison

- ▶ Linear:  $\approx n$  versus Binary Search:  $\log_2(n)$
- ▶ Linear search is *exponentially worse*
- ▶ Binary search is *exponentially faster*

### Perspective

- ▶ Suppose we have a database of 1 trillion,  $10^{12}$  elements
- ▶ Unsorted using linear search:  $\approx 5 \times 10^{11}$  comparisons
- ▶ Sorted (“indexed”) using binary search:  $\approx \log_2(10^{12}) \approx 40$  comparisons

### Another Perspective

Growth Rate

- ▶ Suppose we *double* the input size:  $n \rightarrow 2n$
- ▶ Linear search would require  $n \rightarrow 2n$  comparisons
- ▶ Doubling the input size doubles the number of comparisons
- ▶ Binary search:  $\log_2(n) \rightarrow \log_2(2n)$
- ▶  $\log_2(2n) = \log_2(n) + 1$
- ▶ Doubling the input size only adds one more comparison!

# Part III: Selection Sort

### Introduction

- ▶ To exploit binary search we need to be able to sort
- ▶ Many different sorting algorithms each with different properties
- ▶ Bubble Sort, Selection Sort, Insertion Sort, Quick Sort, Merge Sort, Heap Sort, Tim Sort, etc.
- ▶ Some efficient, some inefficient
- ▶ Start with a simple implementation: Selection Sort

## Basic Idea

- ▶ Search through the array and find the minimal element
- ▶ Swap it with the first element
- ▶ Proceed with the remainder of the array
- ▶ In general:
  - ▶  $i$ -th iteration: find minimal element in `arr[i]` through `arr[n-1]`
  - ▶ Swap it with `arr[i]`
  - ▶ Stop at  $i = n - 1$  (last element is already sorted)
- ▶ Demonstration

## Selection Sort Example

Iteration 1

	swap							
index	0	1	2	3	4	5	6	7
contents	2	42	62	7	20	102	34	47

## Selection Sort Example

Iteration 2

		swap						
index	0	1	2	3	4	5	6	7
contents	2	7	62	42	20	102	34	47

```
1 void selectionSort(int *arr, int n) {
2
3     for(int i=0; i<n-1; i++) {
4         int minIndex = i;
5         for(int j=i+1; j<n; j++) {
6             if(arr[j] < arr[minIndex]) {
7                 minIndex = j;
8             }
9         }
10        //swap
11        int temp = arr[i];
12        arr[i] = arr[minIndex];
13        arr[minIndex] = temp;
14    }
15 }
```

## Analysis

- ▶ Selection sort is simple, but naive and inefficient
- ▶ How bad is it?
- ▶ How many comparisons does selection sort make on an array of size  $n$ ?
  - ▶ First iteration:  $n - 1$  comparisons
  - ▶ Second iteration:  $n - 2$  comparisons
  - ▶  $i$ -th iteration:  $n - i$  comparisons
  - ▶ Last iteration: 1 comparison
  - ▶ In total:

$$1 + 2 + 3 + \dots + (n - 2) + (n - 1) = \frac{n(n-1)}{2} = \frac{1}{2}n^2 + \frac{1}{2}n$$

## Perspective

- ▶ Selection sort is a *quadratic*,  $\approx n^2$  sorting algorithm
- ▶ How bad is this?
- ▶ Sorting the database of 1 trillion,  $10^{12}$  elements requires

$$\approx 5 \times 10^{23}$$

- ▶ 500 “Sextillion” comparisons
- ▶ NVIDIA GTX 1080Ti: 11.3 TeraFLOPS

$$\frac{5 * 10^{23} \text{ operations}}{11.3 * 10^{12} \text{ ops/sec}} = 1,402.157 \text{ years}$$

- ▶ Not feasible for even “moderately large” inputs

## Another Perspective


- ▶ Double the size of the array:  $n \rightarrow 2n$
- ▶ Number of comparisons grows:
$$n^2 \rightarrow (2n)^2 = 4n^2$$
- ▶ Doubling the input *quadruples* the number of operations
- ▶ Four times slower!

- $$n^2 \rightarrow (2n)^2 = 4n^2$$

## Part IV: Quick Sort

## Basic Idea

## Quick Sort Example



## Analysis

- ▶ Best/Worst/Average case analysis
- ▶ Quick Sort makes roughly  $n \log_2(n)$  comparisons
- ▶ *Much* better than  $n^2$
- ▶ Comparisons

## Analysis

- ▶ Sorting the database of 1 trillion,  $10^{12}$  records
- ▶ Comparisons:

$$10^{12} \cdot \log_2 10^{12} \approx 4 \times 10^{13}$$

- ▶ 40 trillion comparisons
- ▶ NVIDIA GTX 1080Ti: 11.3 TeraFLOPS

$$\frac{4 \times 10^{13} \text{ operations}}{11.3 \times 10^{12} \text{ ops/sec}} = 3.5 \text{ seconds}$$

- ▶ Very feasible

## Analysis

- ▶ Consider doubling the input size:  $n \rightarrow 2n$
- ▶ Number of comparisons:

$$n \log_2(n) \rightarrow 2n \log_2(2n)$$

- ▶  $2n \log_2(2n) = 2n \log_2(n) + 2n$
- ▶ Roughly only twice as many
- ▶ Often referred to as *quasilinear*

## Part V: Sorting in Practice

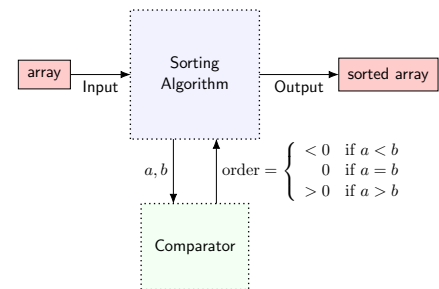
## In Practice

- ▶ Don't "roll your own" searching/sorting algorithms
- ▶ Use standard library functions
- ▶ *But:* we don't want dozens of different functions one for each type of variable or criteria that we want to sort with respect to
- ▶ Want ONE generic solution that can sort *any* type of data by *any* criteria
- ▶ One sorting function to sort them all

## Comparators

- ▶ Solution: use one *generic* sorting function
- ▶ Needs to know how to order two elements,  $a, b$
- ▶ Are they in order or do they need to be swapped?
- ▶ Solution: A *comparator* function
- ▶ Given two elements  $a, b$  it returns:
  - ▶ *something* negative if  $a < b$
  - ▶ zero if  $a = b$
  - ▶ *something* positive if  $a > b$

## Comparator Illustration





## Comparators in C

- ▶ In C, a *comparator function* has the following signature
- ▶ `int cmp(const void *a, const void *b);`
- ▶ `const` means we won't change it, only compare it
- ▶ `void *` is a *generic pointer* that can point to anything
- ▶ Recall: `malloc()`

## Standard Pattern

Standard Pattern:

- ▶ Cast the `void *` to a particular data type
- ▶ Use the data's *state* to determine the proper order
- ▶ Return an integer value that expresses the proper order

## Best Practice

Best Practices:

- ▶ Use descriptive function names
- ▶ Be explicit in your comparisons
- ▶ Avoid “tricks”
- ▶ Reuse comparator functionality when possible

## Examples

- ▶ Write a comparator to order integers in non-decreasing order
- ▶ Write a comparator to order integers in non-increasing order
- ▶ Write a comparator to order `Student` structures by NUID
- ▶ Write a comparator to order `Student` structures by GPA
- ▶ Write a comparator to order `Student` structures by last name/first name

# Part VI: Function Pointers

## Function Pointers

- ▶ Now that we have comparator functions: how do we pass them to a generic sorting function?
- ▶ Easy to pass variables by value or by reference
- ▶ How do we pass a function?
- ▶ We need *function pointers*

## Function Pointers

- ▶ Recall: a *pointer* refers to a memory location
- ▶ What is stored in memory?
- ▶ Variables, arrays, data, *everything*
- ▶ A program's code is stored in memory, including its *functions*
- ▶ We can create pointers that point to memory locations that contain functions!
- ▶ Function pointers allow us to “pass” a function to another function
- ▶ Called “callback” functions
- ▶ Demonstration

## Function Pointers: Demo

```
1 //create a pointer called ptrToFunc that can point to a
2 //function that returns an integer and takes three arguments:
3 //(int, double, char)
4
5 int (*ptrToFunc)(int, double, char) = NULL;
6
7 //declare a pointer that can point to math's sqrt function
8 double (*ptrToSqrt)(double) = NULL;
9
10 //let's make ptrToSqrt point to the sqrt function
11 ptrToSqrt = sqrt;
12
13 //you can call a function via its pointer:
14 double x = ptrToSqrt(2.0);
15
16 //careful: you can reassign standard library functions:
17 sqrt = sin;
18 double y = sqrt(3.14159); //0
19 //don't do this
20
21 //a function that takes another function:
22 void runFunction(double x, double (*func)(double)) {
23     //run func on x:
24     double y = func(x);
25 }
```

## Part VII: Searching & Sorting in C

### Searching & Sorting in C

- ▶ To make generic searching & sorting functions, we need to pass in a comparator
- ▶ Function pointers allow us to do this
- ▶ The array to be searched/sorted is also generic, `void *`
- ▶ Demonstration: generic linear search

### Generic Linear Search

```
1 /**
2  * This function takes an array of integers
3  * and searches it for the given key, returning
4  * the index at which it finds it, or -1 if no
5  * such element exists.
6  */
7 int linearSearch(const void *key, const void *arr, int n, int size, int (*compar) (const void *, const void *)) {
8     for(int i=0; i<n; i++) {
9         if(compar(key, (arr + i * size)) == 0) {
10             return i;
11         }
12     }
13     return -1;
14 }
15 }
```

### Sorting in C

- ▶ The standard C library provides a generic sorting function

```
1 void qsort(void *base,
2            size_t nel,
3            size_t size,
4            int (*compar)(const void *, const void *));
```

- ▶ `base` is the array of elements to be sorted
- ▶ `nel` is the number of elements in the array
- ▶ `size` is the number of bytes each element takes
- ▶ `compar` the comparator function you want to use to order elements
- ▶ Demonstration

## Binary Search in C

- ▶ The standard C library provides a generic binary search function

```
1 void * bsearch(const void *key,  
2             const void *base,  
3             size_t nel,  
4             size_t size,  
5             int (*compar) (const void *, const void *));
```

- ▶ Returns a pointer to the element that “matches” the key  
(an element such that the comparator returns 0)
- ▶ Returns `NULL` if no such element
- ▶ Assumes the array is sorted in the same order as defined by `compar`
- ▶ Demonstration